

# Speeding to resolution with human-in-the-loop automation

Tina Huang, Founder and CTO

# Zoom:

August 24, 2020



# Slack:

May 12, 2020



## Connectivity Issues

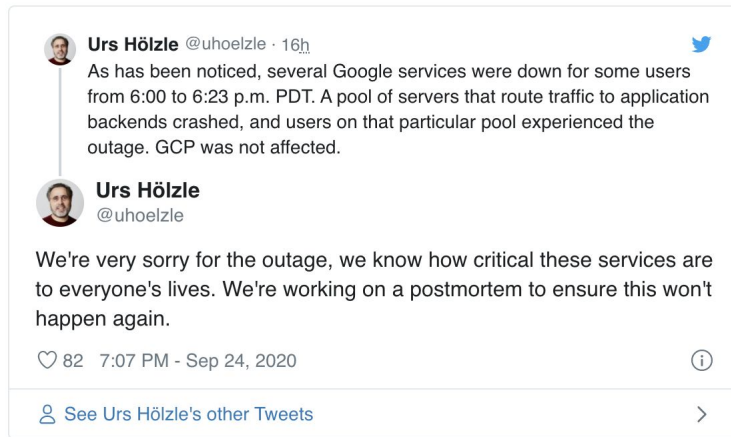
Today at 3:58 PM PDT | [See outage details](#)

# Gmail:

August 20, 2020  
September 24, 2020

# Microsoft 365:

September 28, 2020



A screenshot of a tweet from Urs Hölzle (@uhoelzle) posted 16 hours ago. The tweet contains two paragraphs of text. The first paragraph reports a Google service outage affecting some users from 6:00 to 6:23 p.m. PDT, mentioning that application backends crashed and users on that pool experienced the outage, while GCP was not affected. The second paragraph is an apology, stating that the outage was critical for many users and that a postmortem is being conducted to prevent future occurrences. The tweet has 82 likes and was posted on September 24, 2020, at 7:07 PM. A link to see other tweets by Urs Hölzle is provided at the bottom.

**Urs Hölzle** @uhoelzle · 16h

As has been noticed, several Google services were down for some users from 6:00 to 6:23 p.m. PDT. A pool of servers that route traffic to application backends crashed, and users on that particular pool experienced the outage. GCP was not affected.

**Urs Hölzle** @uhoelzle

We're very sorry for the outage, we know how critical these services are to everyone's lives. We're working on a postmortem to ensure this won't happen again.

♥ 82 7:07 PM - Sep 24, 2020 ⓘ

[See Urs Hölzle's other Tweets](#) ➔



A screenshot of a tweet from Microsoft 365 Status (@MSFT365Status) posted on September 28, 2020, at 3:44 PM. The tweet states that Microsoft is investigating an issue affecting access to multiple Microsoft 365 services and that they are working to identify the full impact and provide more information shortly. The tweet has 3.3K likes and 2.9K people are tweeting about it. An information icon is visible in the bottom right corner.

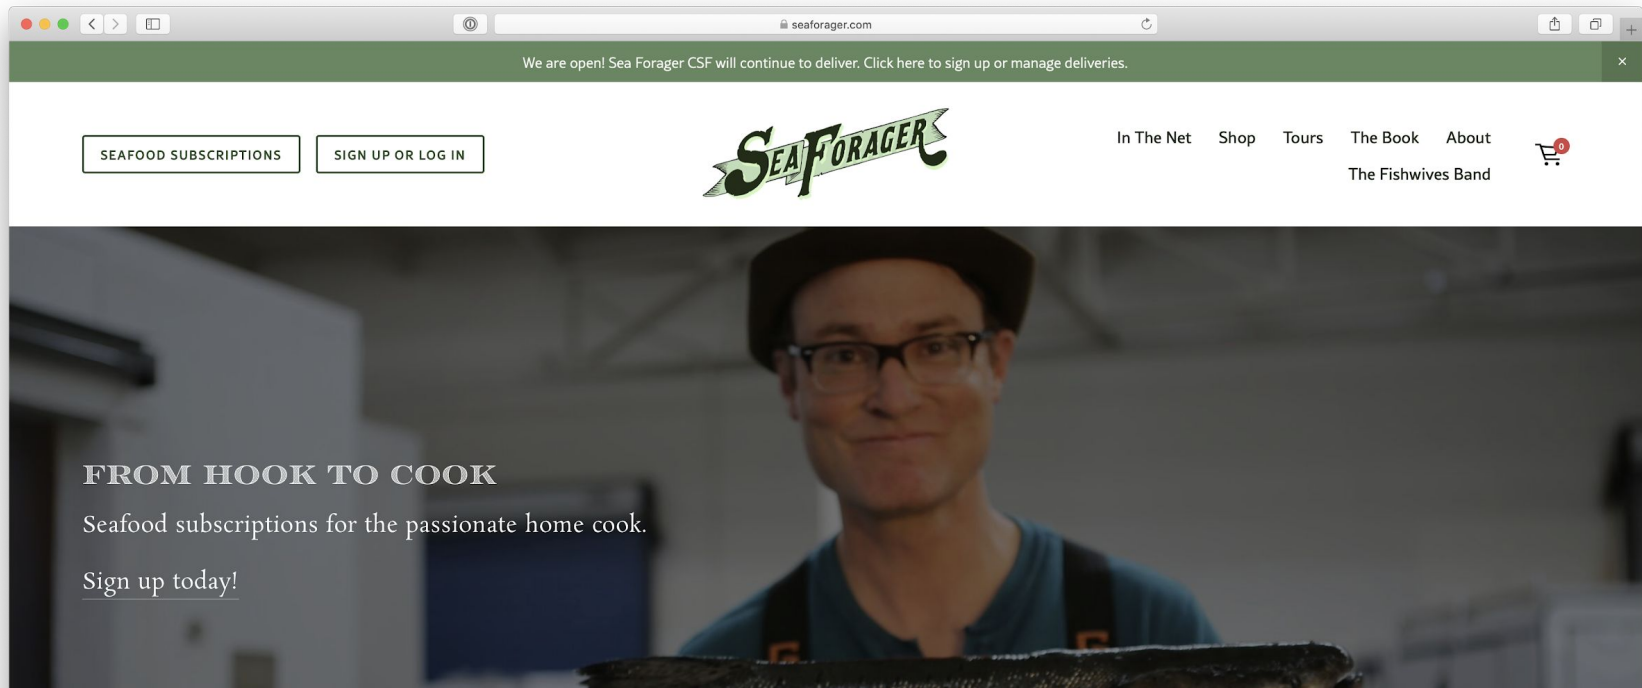
**Microsoft 365 Status** @MSFT365Status

We're investigating an issue affecting access to multiple Microsoft 365 services. We're working to identify the full impact and will provide more information shortly.

3:44 PM · Sep 28, 2020 ⓘ

♥ 3.3K 💬 2.9K people are Tweeting about this

# Digital dependency is at an all time high.




**Agile requires us to  
rethink operations and  
how to make reliability  
sustainable.**



**Agile makes customer support an integral part of reliability.**



**Sustainable agile  
requires faster  
resolution and  
continuous learning**







**Automation and  
post-mortems are not  
the panacea for incident  
management.**



**Resolving faster is more  
than the technical, it's  
about helping people  
communicate better.**



# Why is automation hard?



# Automation is scary



# Nuance is hard to automate



# Innovation breeds incidents



# APIs are hard




# Let's rethink automation





# Human in the loop automation is checklists on steroids




# What makes a great SRE?

- Having lots of operational expertise
- Lots of domain knowledge
- Great operational and organizational intuition




# Turning human intuition into human judgement




**With so much data,  
why is incident  
management not  
more data driven?**



# Incident management is fundamentally broken.




**At 3am no one is  
thinking about capturing  
data**



**You need an incident  
management platform to  
capture data**



# #BeachOps





**Having data isn't enough. How can we make data accessible and useful during incidents?**



**Make data access  
friendly with search and  
recommendations**



# Automating the learning process



Iggy's Surfshop ▾

● Tina Huang



- 🔍 Threads
- @ Mentions & reactions
- 📄 Drafts
- 🔖 Saved items
- 👤 People
- ⋮ Apps
- 📁 Files
- ↑ Show less

▼ Channels



# general

# alerts

# oncall

# random

# TR-1234

+ Add a channel

▼ Direct messages



♥ Slackbot

● Tina Huang (you)

+ Invite people

▼ Apps



#oncall

☆ | 👤 4 | ⚙ 1 | ✎ Add a topic



Q Search



## Web 500s Are Over Threshold



Jump to...

Capacity Problems ▾

## Capacity Problems

- Check the ELB logs to see if there are a lot of 500s there. **503s** indicate that the underlying service is unavailable, **504s** indicate service timeout. In either case, adding capacity probably helps.
- Check is the **ECS metrics** . If CPU or memory are high, we probably need more capacity. Before you add capacity, check that we are actually running the desired number of tasks. If not, sometimes that means a task is wedged and restarting ECS tasks can help. Otherwise, add some capacity.

Restart ECS

ECS Metrics

Scale ECS

Trigger Terraform using Jenkins

Close

Transp  
joined #Transp  
JIRA iss  
IS-963  
The pay

Tasks

Current Runbook

Browse Runbooks

Message #oncall



B

I



&lt;/&gt;



Aa



Iggy's Surfshop ▾

● Tina Huang



- 🔍 Threads
- @ Mentions & reactions
- 📄 Drafts
- 🔖 Saved items
- 👤 People
- ⋮ Apps
- 📁 Files
- ↑ Show less

▼ Channels



# general

# alerts

# oncall

# random

# TR-1234

+ Add a channel

▼ Direct messages



♥ Slackbot

● Tina Huang (you)

+ Invite people

▼ Apps



#oncall

☆ | 👤 4 | ⚙ 1 | ✎ Add a topic



Q Search



## Web 500s Are Over Threshold



Jump to...

Capacity Problems ▾

## Scale ECS

## Service

Execution ▾

## Instances

6 ▾

Cancel

Run

Transparency  
joined #Transparency  
JIRA issue created by @and

## IS-963 Payments server is down - Priority P1

The payments server was erroring and was offline at 3:06 PM today, as reported. Ongoing assessment in TR-1234.

Tasks

Current Runbook

Browse Runbooks

Message #oncall



Aa



Iggy's Surfshop ▾

● Tina Huang



- 🔍 Threads
- @ Mentions & reactions
- 📄 Drafts
- 🔖 Saved items
- 👤 People
- ⋮ Apps
- 📁 Files
- ↑ Show less

▼ Channels



# general

# alerts

# oncall

# random

# TR-1234

+ Add a channel

▼ Direct messages



♥ Slackbot

● Tina Huang (you)

+ Invite people

▼ Apps



#oncall

☆ | 👤 4 | 🌟 1 | ✎ Add a topic



Q Search

**Transposit** APP 3:10 PM

joined #tr-1234 along with @tina

**Transposit** APP 3:10 PM

JIRA issue created by @tina

**IS-963 Payments server is down - Priority P1**

The payments server was erroring and was offline at 3:06 PM today, as reported. Ongoing assessment in TR-1234.

Tasks

Current Runbook

Browse Runbooks

@Tina Huang ran Scale ECS

Restarting tasks for ecs service excution on cluster production-cluster. Updates will be posted in this channel.

**12:39:54 PM (PDT):** (service excution) has started 6 tasks: (task 5154eabb-f324-4c38-88f9-00fe57991692) (task 747a8acc-f076-407e-8177-a3cf39247b50)**12:39:55 PM (PDT):** (service excution) registered 6 targets in (target-group arn:aws:elasticloadbalancing:us-west-2:668521361712:targetgroup/production-cluster-excution/6c7c5b650e48c849)**12:39:56 PM (PDT):** (service excution) deregistered 6 targets in (target-group arn:aws:elasticloadbalancing:us-west-2:668521361712:targetgroup/production-cluster-excution/6c7c5b650e48c849)**12:39:57 PM (PDT):** (service excution) has begun draining connections on 1 tasks**12:39:57 PM (PDT):** (service excution) has stopped 6 running tasks: (task ef7aa378-ed05-4cdf-bf10-bb10b7a49f1e) (task ebd20d4e-1be5-4862-9d2c-9985dc46906b)**12:39:59 PM (PDT):** (service excution) has reached a steady state. :white\_check\_mark:

Message #oncall



Aa





Show All Events

T-132

T-133

T-134

T-135

T-136

T-137

Alerts

PagerDuty

OpsGenie

Zendesk

Time ↓

Event



Event Type

People

Date

Time

12:43 PM **Leslie Bauman** @leslie

"Looks like someone on the analytics team is running queries against prod. I'm going to kill that one.."

...

12:43 PM **Running DB Queries**

run by Leslie Bauman

...

```
SELECT * FROM users AS U JOIN teams AS T WHERE T.id = U.teamId AND T.id IN  
(SELECT * FROM teams WHERE org = "Engineering") (6009ms)
```

```
SELECT * FROM teams WHERE id = '529' (29ms)
```

12:43 PM **Running DB Queries**

run by Leslie Bauman

...

12:43 PM **Nina Yang** @nina

"Not sure why an alert didn't fire for the DB, but CPU load seems high"

...

12:43 PM **Nina Yang** @nina

"Looks like the website is lagging when customers add items to the shopping cart"

...

12:43 PM **Leslie Bauman** @leslie

@nina I'm not really sure what's going on but can you help poke around?

...

12:43 PM **Nina Yang** @nina

Started runbook Latency Exceeds Maximum

...

12:43 PM **Create a new JIRA issue** (Done)**Nina Yang** created a new Jira issue created  
[Jira T-539](#)

...

12:43 PM **Alert**

Pagerduty: Web 500s are over threshold

...

Runbooks



Show All Events

~ T-132

~ T-133

~ T-134

~ T-135

~ T-136

~ T-137

~ Alerts

~ PagerDuty

~ OpsGenie

~ Zendesk

Time ↓ Event



12:43 PM **Nina Yang** @nina  
Started runbook Latency Exceeds Maximum



12:43 PM **Alert**  
Pagerduty: Web 500s are over threshold



Event Type

People

Date

Time

## Latency Exceeds Maximum

Jump to...  
▼

## Capacity Problems

- Check the ELB logs to see if there are a lot of 500s there. **503s** indicate that the underlying service is unavailable, **504s** indicate service timeout. In either case, adding capacity probably helps.
- Check the **ECS metrics**. If CPU or memory are high, we probably need more capacity. Before you add capacity, check that we are actually running the desired number of tasks. If not, sometimes that means a task is wedged and restarting ECS tasks can help. Otherwise, add some capacity.

Restart ECS

ECS Metrics

Trigger Terraform

Scale ECS

Commonly used following alert  
Pagerduty: Web 500s are over threshold

Next step: Server Errors

Previous

Next





Show All Events

T-132

T-133

T-134

T-135

T-136

T-137

Alerts

PagerDuty

OpsGenie

Zendesk

Time ↓

Event



Event Type

People

Date

Time



12:43 PM Leslie Bauman @leslie

"Looks like someone on the analytics team is running queries against prod"



12:43 PM Running DB Queries

run by Leslie Bauman

```
SELECT * FROM users AS U JOIN teams AS T WHERE T.id = U.teamId AND  
(SELECT * FROM teams WHERE org = "Engineering") (6009ms)
```

```
SELECT * FROM teams WHERE id = '529' (29ms)
```



12:43 PM Running DB Queries

run by Leslie Bauman



12:43 PM Nina Yang @nina

"Not sure why an alert didn't fire for the DB, but CPU load seems high"



12:43 PM Nina Yang @nina

"Looks like the website is lagging when customers add items to the shopping cart"



12:43 PM Leslie Bauman @leslie

@nina I'm not really sure what's going on but can you help poke around?



12:43 PM Nina Yang @nina

Started runbook Latency Exceeds Maximum



12:43 PM Create a new JIRA issue (Done)

Nina Yang created a new Jira issue created  
[Jira T-539](#)



12:43 PM Alert

Pagerduty: Web 500s are over threshold

Event type

Workflow

Within alert type

Web 500s are over threshold

Clear

Apply

Runbooks



Show All Events

T-132

T-133

T-134

T-135

T-136

T-137

Alerts

PagerDuty

OpsGenie

Zendesk

Time ↓ Event



Event Type

People

Date

Time



12:43 PM **Running DB Queries**  
run by Leslie Bauman



```
SELECT * FROM users AS U JOIN teams AS T WHERE T.id = U.teamId AND T.id IN  
(SELECT * FROM teams WHERE org = "Engineering") (6009ms)
```

```
SELECT * FROM teams WHERE id = '529' (29ms)
```

T-82



8:33 AM **Create a new JIRA issue** (Done)  
Nina Yang created a new Jira issue created  
[Jira T-539](#)



T-123

June 8, 2020



12:43 PM **Scale ECS**  
run by Leslie Bauman



```
12:39:54 PM (PDT): (service excution) has started 6 tasks: (task  
5154eabb-f324-4c38-88f9-00fe57991692) (task  
747a8acc-f076-407e-8177-a3cf39247b50)  
12:39:55 PM (PDT): (service excution) registered 6 targets in (target-group  
arn:aws:elasticloadbalancing:us-west-2:668521361712:targetgroup/production-cl  
uster-excution/6c7c5b650e48c849)  
12:39:56 PM (PDT): (service excution) deregistered 6 targets in (target-group  
arn:aws:elasticloadbalancing:us-west-2:668521361712:targetgroup/production-cl  
uster-excution/6c7c5b650e48c849)
```

T-142

Runbooks



## Development



## Auths &amp; settings

Authorize data connections



## Stash



## Dynamic Tasks

## Operations

list\_commits  
SQLinput\_prompt  
PY

## Data connections

GitHub  
transposit/github

Run

Discard code

Commit code

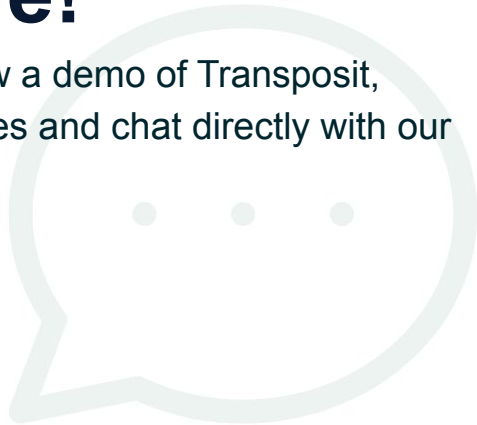
```
1 SELECT * FROM github.list_commits
2 WHERE owner='tina@transposit.com'
3 AND repo='transposit/terraform'
5 AND author=*
6 AND since='09/01/2020'
7 LIMIT 100
8
9
10
11
```

# **The Future of Operations is Now**



# Join me in our Slack Q&A to chat more!

Visit our booth to view a demo of Transposit,  
enter to win cool prizes and chat directly with our  
team!



Join us for

